

BRBON

V 0.3.1

A Binary Object Notation by Balancing Rock.

Introduction

BRBON is a binary storage format specification. It started out as a binary version for JSON but the requirement for speed has rendered some JSON aspects obsolete. Still, the JSON origins can be recognised in the object oriented approach.

The major design driver for BRBON is speed. Speed in access and speed in loading & storage. A secondary goal was predictable speed. BRBON is designed for use in a web server application and thus the access times should be constant and not fluctuate widely.

Strictly speaking BRBON only refers to the format specification, not an implementation. However a reference API is offered for free under the same name.

- Vectored access: The format includes vectors that can be used to quickly traverse the entire structure (in both directions).
- Named objects: Objects may be named for identification, a hash value is included for faster name recognition.
- Aligned to 8 byte boundaries: Most of the elements in the structure are aligned to 8 byte boundaries. This ensures alignment for all types.
- All standard types are supported: Bool, Int8, Int16, Int32, Int64, UInt8, UInt16, UInt32, UInt64, Float32, Float64, String, Binary, UUID
- Collection types: Array, Dictionary, Sequence, Table.
- Special type: Null.
- Array's are packed for optimal storage density.
- The maximum Item capacity is 2GB (Int32.max). (The unfinished Block specification will include the capability for larger structures)

At this point in time BRBON is considered incomplete. The 'Item' specification has been completed but the 'Block' specification is still open. The Item specification can be used for local storage needs and is useful on its own. The reference API has been prepared for the Block extension by taking the endianness of the data into account. Endianness will be part of the block specification.

Two alternatives for BRBON have been considered: BSON and BinSON. However both were found lacking with respect to the capabilities for high speed (memory mapped) access.

History

Date	Version	Changes
2018-02-24	V0.2	First Draft
2018-03-22	V0.3	Restyling, removed count from count/Value and renamed it small-value. No longer backwards compatible to V0.2
2018-03-30	V0.3.1	Added the UUID type

Block

A Block is a container of Items. It consists of a Block-Header, followed by Items, followed by a Block-Footer.

The purpose of a block is:

- Provide intra-system information (a.o. endianness)
- Allow sizes > 2GB
- Synchronisation of BRBON data between systems
- Allow detection of data inconsistency
- Allow encryption

TBD.

Item

The purpose of an Item is to describe the data contained in its payload as well as to provide path-based access to the payload. An Item can contain an hierarchy of other Items within it.

Item =

```
Item-Header,      (* 4 bytes *)
Item-Byte-Count, (* 4 bytes, value is always a multiple of 8 *)
Parent-Offset,    (* 4 bytes *)
Small-Value       (* 4 bytes, Payload if fixed and fits in 32 bits *)
[,Name-Field]     (* N x 8 bytes *)
[,Value-Field]    (* Payload if variable or fixed and > 32 bits *)
[,Filler]         (* Filler bytes to round out the Item-Byte-Count *)
```

Rationale: The item header contains all data needed to know what is where in the item. The item byte count contains the number of bytes of the entire item. This allows skipping over items at high speed. The name field (if present) is always located at the same offset from the first byte. This allows fast name-checking of a series of successive items.

Item Header

The item header defines the type of payload, its options and the possible presence of a name-field. The item header is made of 4 bytes, each with its own definition.

Item-Header =

```
Item-Type,          (* 1 byte *)
Item-Options,       (* 1 byte, set to 0 if unused *)
Item-Flags,         (* 1 byte *)
Name-Field-Byte-Count (* 1 byte *)
```

Rationale: This is essential data that most APIs will need to access first when looking up an item.

Item-Type

The item-type is a byte field specifying the type of payload in either the small-value field or the value-field.

The following types are defined:

Hex code	Type description	Location
0	Illegal value.	n.a.
0x01	Null	n.a.
0x02	Bool	Small-Value
0x03	Int8	Small-Value
0x04	Int16	Small-Value
0x05	Int32	Small-Value
0x06	Int64	Value-Field
0x07	UInt8	Small-Value

Hex code	Type description	Location
0x08	UInt16	Small-Value
0x09	UInt32	Small-Value
0x0A	UInt64	Value-Field
0x0B	Float32	Small-Value
0x0C	Float64	Value-Field
0x0D	String	Value-Field
0x0E	CRC String	Value-Field
0x0F	Binary	Value-Field
0x10	CRC Binary	Value-Field
0x11	Array	Value-Field
0x12	Dictionary	Value-Field
0x13	Sequence	Value-Field
0x14	Table	Value-Field
0x15	UUID	Value-Field
0x16-0x7F	Reserved (do not use)	n.a.
0x80-0xFF	Available for user definitions (free for all).	n.a.

For more information on the types and how they are coded, see the Types section.

Item-Options

A field of 8 bits witch is currently unused and must be zero. If there is a non-zero bit, it must be assumed that the data of this item cannot be interpreted.

Item-Flags

The flag area is for used during runtime. When stored or transferred the bits are undefined.

Name-Field-Byte-Count

Contains an UInt8 value indicating the number of bytes in the name field. This value is/must always be a multiple of 8. Its maximum value is 248.

When the name field is not present, this value is zero.

The least significant 3 bits are reserved and should be zero. If one of these bits is set, it is likely an error in the data structure. It is recommended to raise an error and/or disregard the item.

Item Byte Count

The number of bytes the item occupies.

Range: 16 ... Int32.max.

It is always a multiple of 8, the lower 3 bits must be zero.

The value 0 is a special case, it may point to self if self is the root item or it may point to the root item if self is contained in the root item. To differentiate between them, compare the address of the 'parent' with the address of 'self', if it is equal, the item is the root item.

Rationale: Creating an explicit size definition at the top level allows for faster hierarchy traversing. API's do not need to evaluate an item structure in order to find a successive item.

Parent Offset

The offset of the parent item relative to the first item byte in a block. Items can be contained within items for the array, dictionary, sequence and table item types. If an item is contained in another item, the parent offset allows an API to traverse to the parent item.

Note that this value can be recalculated at any time.

Rationale: APIs frequently must be able to traverse up and down in an object hierarchy. Since BRBON species a memory mapped structure all the data necessary for this must be present in the BRBON structure specification itself. Even if there is no implicit need to store or transfer this data.

Small Value

If an item-type with a fixed size is small enough to fit in a 32 bit field, this field will be used for storage.

Rationale: This allows the smallest possible item size while at the same time allowing for variable length names. Alternatives have been tried, but always resulted in either minimum item sizes of 24 bits or slower access speeds. While item size is not a main driver, the alternatives that use a minimum of 24 bits were not faster and are thus not preferable.

Name Field

Name-Field =

```
Name-CRC,          (* 2 bytes, CRC16, ARC, rev polynomial 0x8005 *)
UTF8-Byte-Count, (* 1 byte, the number of used UTF8 bytes *)
UTF8-Byte-Code    (* UTF8-Byte-Code bytes *)
[, Filler]        (* Filler, optional *)
```

Name-CRC: The hash value is a CRC16 using the reverse polynomial 0x8005 (= 0xA001 reversed) and 0 as the start seed.

UTF8-Byte-Count: This is the number of bytes used by the UTF8 byte code that constitutes the name.

UTF8-Byte-Code: The UTF-8 byte code that can be converted into a valid string.

Filler: The filler is a number of bytes that ensure that the number of bytes in the name field is a multiple of 8. This limits the total name field to a maximum of 248 bytes.

Note: Subtracting the 3 bytes used by CRC and UTF8 byte count gives a maximum size of 245 bytes for the name characters.

Rationale: The CRC allows for faster finding of items by name. If the CRC does not match, the UTF8 byte code does not need to be compared. Names longer than 245 characters seem unlikely, they would make for unreadable code.

Value Field

The layout of the value field is given in the section Types.

Filler

Filler data must be present if the value field is not a multiple of 8 bytes. It then rounds of the item byte count to a multiple of 8 bytes. More often though the filler may be present because a designer wants to ensure that some area is available for later expansion without having to shift the entire data field following this item.

Types

Null

The null is a special type that has no associated value. It is intended to be used as a placeholder that can be changed into any other type.

It can contain filler data.

A null cannot be stored in an array or table.

Bool

The value field is a single byte containing either zero (false) or non-zero (true).

When stored in an item the value will be stored in the first byte of the small-value field.

When stored in an array or table the value will be stored in a single byte, the first byte in an element or field.

Int8

A single byte, range -128 ... +127. (0x80 ... 0x7F)

When stored in an item the value will be stored in the first byte of the small-value field.

When stored in an array or table the value will be stored in a single byte, the first byte in an element or field.

Int16

A 2-byte value, range -32,768 ... +32,767 (0x8000 ... 0x7FFF)

When stored in an item the value will be stored in the first two bytes of the small-value field.

When stored in an array or table the value will be stored in the first 2 consecutive bytes in an element or field.

Endianness is specified by the block endianness or when the block is absent the machine endianness.

Int32

A 4-byte value, range -2,147,483,648 ... +2,147,483,647 (0x8000_0000 ... 0x7FFF_FFFF)

When stored in an item the value will be stored in the small-value field.

When stored in an array or table the value will be stored in the first 4 consecutive bytes in an element or field.

Endianness is specified by the block endianness or when the block is absent the machine endianness.

Int64

An 8-byte value, range -9,223,372,036,854,775,808 ... +9,223,372,036,854,775,807 (0x8000_0000_0000_0000 ... 0x7FFF_FFFF_FFFF_FFFF)

When stored in an item the value will be stored in the first eight consecutive bytes in the value field.

When stored in an array or table the value will be stored in the first eight consecutive bytes in the value field.

The small-value field is not used.

Endianness is specified by the block endianness or when the block is absent the machine endianness.

UInt8

A single byte, range 0... +255. (0x00 ... 0xFF)

When stored in an item the value will be stored in the first byte of the small-value field.

When stored in an array or table the value will be stored in a single byte, the first byte in an element or field.

UInt16

A 2-byte value, range 0 ... +65,535 (0x0000 ... 0xFFFF)

When stored in an item the value will be stored in the first two bytes of the small-value field.

When stored in an array or table the value will be stored in the first 2 consecutive bytes in an element or field.

Endianness is specified by the block endianness or when the block is absent the machine endianness.

UInt32

A 4-byte value, range 0 ... +4,294,967,295 (0x0000_0000 ... 0x8FFF_FFFF)

When stored in an item the value will be stored in the small-value field.

When stored in an array or table the value will be stored in the first 4 consecutive bytes in an element or field.

Endianness is specified by the block endianness or when the block is absent the machine endianness.

UInt64

An 8-byte value, range 0 ... +18,446,744,073,709,551,615 (0x8000_0000_0000_0000 ... 0x7FFF_FFFF_FFFF_FFFF)

When stored in an item the value will be stored in the first eight consecutive bytes in the value field.

When stored in an array or table the value will be stored in the first eight consecutive bytes in the value field.

The small-value field is not used.

Endianness is specified by the block endianness or when the block is absent the machine endianness.

Float32

A 4-byte value. Accurate to about 6 decimals, range approx 1.1e-38 to 3.4e38

When stored in an item the value will be stored in the small-value field.

When stored in an array or table the value will be stored in the first 4 consecutive bytes in an element or field.

Endianness is specified by the block endianness or when the block is absent the machine endianness.

Float64

An 8-byte value, range approx 2.2e-308 to 1.7e+308

When stored in an item the value will be stored in the first eight consecutive bytes in the value field.

When stored in an array or table the value will be stored in the first eight consecutive bytes in the value field.

The small-value field is not used.

Endianness is specified by the block endianness or when the block is absent the machine endianness.

String

`String = byte-count [{, UTF8-Byte-Code}]`

The string is a sequence of bytes representing the UTF8 code of the string. The UTF8 byte code is preceded by a count of the number of bytes in the byte code.

The layout is as follows:

Offset	Content	Units	Bytes
0	Byte Count	UInt32	4
4	UTF8 Byte Code	Bytes	Byte Count
4 + Byte Count	Filler	Bytes	As needed

This layout is either stored in the Value-Field of an item, the element of an array or the field in a table.

Endianness is specified by the block endianness or when the block is absent the machine endianness.

CRC String

`CRC-String = CRC32, Byte-Count [{, UTF8-Byte-Code}]`

The string is a sequence of bytes representing the UTF8 code of the string. The UTF8 byte code is preceded by the CRC32 of the bytes in the UTF8 byte code and a count of the number of bytes in the byte code.

The CRC32 uses a seed value of 0.

The layout is as follows:

Offset	Content	Units	Bytes
0	CRC32	UInt32	4
4	Byte Count	UInt32	4
8	UTF8 Byte Code	Bytes	Byte Count
8 + Byte Count	Filler	Bytes	As needed

This layout is either stored in the Value-Field of an item, the element of an array or the field in a table.

Endianness is specified by the block endianness or when the block is absent the machine endianness.

Rationale: This storage structure makes it possible to always use the first 8 bytes of the crc-string for a quick check to see if the following string is **not** equal to a search string. Only when the first 8 bytes are the same, the remaining bytes of the string need to be verified.

Binary

`Binary = byte-count [{, Byte}]`

A binary is a sequence of bytes. The byte sequence is preceded by a count of the number of bytes.

The layout is as follows:

Offset	Content	Units	Bytes
0	Byte Count	UInt32	4

Offset	Content	Units	Bytes
4	Binary	Bytes	Byte Count
4 + Byte Count	Filler	Bytes	As needed

This layout is either stored in the Value-Field of an item, the element of an array or the field in a table.

Endianness is specified by the block endianness or when the block is absent the machine endianness.

CRC Binary

CRC-Binary = CRC32, Byte-Count [{, Byte}]

A binary is a sequence of bytes. The byte sequence is preceded by a CRC32 over the bytes in the binary and a count of the number of bytes.

The CRC32 uses a seed value of 0.

The layout is as follows:

Offset	Content	Units	Bytes
0	CRC32	UInt32	4
4	Byte Count	UInt32	4
8	Binary	Bytes	Byte Count
8 + Byte Count	Filler	Bytes	As needed

This layout is either stored in the Value-Field of an item, the element of an array or the field in a table.

Endianness is specified by the block endianness or when the block is absent the machine endianness.

Rationale: This storage structure makes it possible to always use the first 8 bytes of the crc-binary for a quick check to see if the following binary is **not** equal to a search binary. Only when the first 8 bytes are the same, the remaining bytes of the binary need to be verified.

Array

Contains a sequence of elements. All of the same type and byte count (length). Addressing of the elements is always through their index.

```
Array =
    Reserved,                (* 4 Bytes, must be zero *)
    Element-Type,            (* 1 Byte, see Item Type *)
    Reserved,                (* 3 Bytes, must be zero *)
    Element-Count,          (* The number of elements *)
    Element-Byte-Count      (* The number of bytes in an element *)
    [{, Element }]
```

The layout is as follows:

Offset	Content	Units	Bytes
0	Reserved	UInt32	4
4	Element-Type	Item-Type	1

Offset	Content	Units	Bytes
5	Zero	Bytes	3
8	Element-Count	UInt32	4
12	Element-Byte-Count	UInt32	4
16	Array	Element	Element-Count x Element-Byte-Count
16 + Element-Count x Element-Byte-Count	Filler	Bytes	As needed

When an array is contained in an array or table, it is included as an item.

Sequence

Contains a sequence of other items.

```
Sequence = Count [{,Item}]
```

A sequence does not place any restrictions on the items it contains. (Unlike a dictionary)

The layout is as follows:

Offset	Content	Units	Bytes
0	Reserved	UInt32	4
4	Count	UInt32	4
8	Items	Item	As needed
As needed	Filler	Bytes	As needed

When a sequence is contained in an array or table, is is included as an item.

Dictionary

See Sequence, but each item in the dictionary is guaranteed to have a unique name.

Table

A table is conceptually an array with the same type of dictionary for each element. The key's of the dictionary fields are only present once.

Table =

```

RowCount,                (* 4 bytes, UInt32 *)
ColumnCount,             (* 4 bytes, UInt32 *)
RowsOffset,              (* 4 bytes, UInt32 *)
RowByteCount             (* 4 bytes, UInt32 *)
[{, ColumnDescriptor }] (* 16 bytes * number of columns *)
[{, ColumnName }]        (* N bytes, 8 byte boundary, * num of cols *)
[{, Row }]               (* N bytes, 8 byte boundary, * num of cols *)
```

ColumnDescriptor =

```
ColumnNameCrc16,        (* 2 bytes *)
```

```

ColumnNameByteCount, (* 1 bytes *)
ColumnFieldType,      (* 1 bytes *)
ColumnNameUtf8Offset, (* 4 bytes *)
ColumnFieldOffset,    (* 4 bytes *)
ColumnFieldByteCount  (* 4 bytes *)

```

ColumnName = UTF8-Byte-Count {, UTF8-Byte-Code } (* Max 248 bytes *)

Row = {, ColumnValue } (* 8 byte boundary aligned *)

RowCount: The number of rows in this table (UInt32)

ColumnCount: The number of columns in this table (UInt32)

RowsOffset: The offset from the start of the value field to the first byte of value of the first column in the first row.

RowByteCount: The number of bytes in a row, is equal to the value of all columnValueByteCounts added together.

ColumnDescriptor: There is a column descriptor for each column, containing descriptions of the name and value for that column. Each column descriptor is 16 bytes, and all column descriptors are arranged sequentially.

ColumnNameCrc16: The CRC16 value of the UTF8-byte code that is the name of the column. Using the reverse polynomial 0x8005 (= 0xA001 reversed) and 0 as the start seed.

ColumnNameByteCount: The number of bytes for the UTF8-byte sequence of the column name. (Not the number of used bytes, this is part of the ColumnName field)

ColumnValueType: The type of item stored in this column.

ColumnNameUtf8Offset: The offset from the start of the value area to the sequence of UTF8 byte codes that specify the name of the column.

ColumnValueOffset: The offset from the start of the row where this column's value is located.

ColumnValueByteCount: The number of bytes in the column's value.

The ColumnDescriptor and ColumnName are located in two different area's. The descriptor contains an offset into ColumnNames field. For each descriptor there must be one column name and vice versa.

Column names must be unique inside a table. Column names have a maximum length of 245 UTF8 byte code bytes.

Note: An API should probably set the initial content of a column field when a row is added to all zero to prevent/detect illegal content.

The layout of the table value field is as follows:

Offset	Content	Units	Bytes
0	Row Count	UInt32	4
4	Column Count	UInt32	4
8	Rows Offset	UInt32	4
12	Row Byte Count	UInt32	4
16	Column Descriptors	Column Descriptor	16 x Column Count
16 + 16 x Column Count	Column Names	Column Name	As needed

Offset	Content	Units	Bytes
Rows Offset	Rows (Fields)	Rows	Row Count x Row Byte Count
Rows Offset + Row Count x Row Byte Count	Filler	Bytes	As needed

The layout of the Column Descriptors is as follows:

Offset	Content	Units	Bytes
0	Column Name CRC16	UInt16	2
2	Column Name Field Byte Count	UInt8	1
3	Column Field Type	Item Type	1
4	Column Name UTF8 Offset	UInt32	4
8	Column Field Offset	UInt32	4
12	Column Field Byte Count	UInt32	4

The layout of the Column Names is as follows:

Offset	Content	Units	Bytes
0	Column Name Byte Count	UInt8	1
4	Column Name UTF8 Byte Code	Bytes	Column Name Byte Count

UUID

An array of 16 bytes.

UUID = Bytes (* 16 bytes, UInt8 *)

When stored in an item the value will be stored in the first 16 consecutive bytes in the value field.

When stored in an array or table the value will be stored in the first 16 consecutive bytes in the value field.

The small-value field is not used.

There is no endianness associated with this value.

END OF DOCUMENT